

Performance Metrics Roadmap

1 October 2025, Submission Day

List of Metrics

All refer to an individual classifier-language pair unless otherwise noted

- (Overall) accuracy: percent of correct predictions
- Error type metrics:
 - Illustrating error types (columns)
 - Predicted $\rightarrow$ Actual
 - Count
 - % of all errors
 - (Occurrence Rate [of gold] in Dataset) *This can maybe be omitted; B talked about this, but I honestly don't recall. Unless I remember otherwise, I'll just leave it.*
 - Feature-error associations (columns)
 - Error Type (i.e. $x \rightarrow y$)
 - Feature Values¹, X (number/percent? of tokens² annotated with this feature)
 - Feature Values, Y
 - % of $x \rightarrow y$ errors with this feature bundle
 - % of x-annotated tokens with this feature bundle
 - % of y-annotated tokens with this feature bundle
- Feature performance metrics
 - Baseline (**important**: clarify “baseline” versus “deviation from baseline” and “easier/harder than baseline”)
 - FER (feature error weight)
 - WS (weighted score)
 - **Note**: I currently have four tables in each language section:
 - RFC WS by feature
 - MLP WS by feature
 - RFC FER%, WS, RNS
 - MLP FER%, WS, RNS

Barend pointed out that RNS isn't very useful, and I'm not sure if FER% is either. If FER% isn't useful, I'll just delete these last two tables; if FER% is useful, I'll keep the tables but remove the RNS column.
- Statistical metrics
 - χ^2
 - p -value
 - ϕ -value

¹ Note: For this and the next, I'm not giving all features, only the ones that I've identified as salient (such as 'all +Res, 7/9 +Acc').

² Question: For this and the next, I didn't think to ask whether this should be all tokens in the train+test set or only in the test set; my intuition says test set only.